

TASK 8: INTEGRATING STREAMLIT FRONTEND WITH FLASK REST API

Comprehensive Technical Report: Full Source Code, UI Screenshots, Verification Results & Observations

Course / Platform: L&T Edutech Deep Learning Deployment	Task Identifier: Task 8 - Flask & Streamlit Integration
Frontend Framework: Streamlit 1.30 (Multi-Page Navigation)	Backend Framework: Flask 3.1 & PyTorch 2.x Deep Learning Engine
Integration Test Status: 11 / 11 Test Cases PASSED (100%)	Target Model: DeepHealthRiskNet (3-Layer Neural Network)

1. Executive Summary & Objective

The primary objective of Task 8 is to establish seamless communication between frontend and backend components for real-time deep learning inference. The system connects a multi-page interactive Streamlit user interface (frontend) with a Flask REST API (backend) serving a trained PyTorch 3-layer Artificial Neural Network (DeepHealthRiskNet).

Key Deliverables Achieved:

- Integrated Application: Streamlit Multi-Page UI connected via HTTP requests to Flask REST API endpoints.
- Comprehensive REST API Service: Endpoints for GET /, GET /health, and POST /predict with full JSON payload handling.
- Complete Automated Test Suite: 11 integration test cases validating payload formats, error handling, and latency.
- Full Documentation Package: Complete source code listing, screenshot placeholders, quantitative results, and analytical observations.

2. System Architecture & REST API Endpoints Specification

The client-server architecture decouples user interface logic from neural network inference execution. The Streamlit UI collects patient parameters, formats structured HTTP request payloads, and transmits them to the Flask REST API. The Flask backend validates inputs, executes PyTorch forward pass inference, and returns JSON-encoded predictions and confidence scores.

HTTP Method	Endpoint Path	Description	Expected Payload Format
GET	/	API Root Metadata & Information	None
GET	/health	Health & Readiness Status Probe	None

POST	/predict	Single Patient Real-Time Inference	{"features": [v1, ..., v14]}
POST	/predict	Cohort Batch Ingestion Inference	{"instances": [[...], [...]]}

3. Full Source Code Implementation

Below is the complete, un-truncated source code for all components of the Task 8 integrated application:

3.1 Flask REST API Backend Server (flask_api.py)

Source Code: flask_api.py

```

"""
flask_api.py
-----
Flask REST API Backend for serving PyTorch Deep Learning Model (DeepHealthRiskNet).
Exposes RESTful HTTP endpoints for real-time single and batch predictions, health probes,
and interactive documentation.
"""

import time
import logging
from datetime import datetime, timezone
from flask import Flask, request, jsonify
from model_loader import get_model_engine

# 1. Configure Logging
logging.basicConfig(
    level=logging.INFO,
    format="%(asctime)s [%(levelname)s] %(name)s: %(message)s"
)
logger = logging.getLogger("Flask-DL-Backend")

# 2. Initialize Flask App
app = Flask(__name__)
app.config["JSON_SORT_KEYS"] = False
app.json.compact = False

# 3. Pre-load PyTorch Inference Engine on App Launch
try:
    model_engine = get_model_engine()
    logger.info("PyTorch Model Engine loaded successfully on Flask startup.")
except Exception as e:
    logger.error(f"Failed to initialize PyTorch Model Engine: {str(e)}")
    model_engine = None

def get_iso_timestamp():
    return datetime.now(timezone.utc).isoformat()

def make_error_response(status_code, error_code, message):
    return jsonify({
        "status": "error",
        "error_code": error_code,
        "message": message,
        "timestamp": get_iso_timestamp()
    }), status_code

# -----
# HTTP ERROR HANDLERS
# -----

```

```

@app.errorhandler(400)
def bad_request_error(e):
    return make_error_response(400, "BAD_REQUEST", str(e.description if hasattr(e, 'description') else str(e)))

@app.errorhandler(404)
def not_found_error(e):
    return make_error_response(404, "NOT_FOUND", f"Endpoint '{request.path}' not found.")

@app.errorhandler(405)
def method_not_allowed_error(e):
    return make_error_response(405, "METHOD_NOT_ALLOWED", f"Method '{request.method}' not allowed on '{request.path}'.")

@app.errorhandler(415)
def unsupported_media_type_error(e):
    return make_error_response(415, "UNSUPPORTED_MEDIA_TYPE", "Content-Type must be 'application/json'.")

@app.errorhandler(422)
def unprocessable_entity_error(e):
    return make_error_response(422, "UNPROCESSABLE_ENTITY", str(e.description if hasattr(e, 'description') else str(e)))

@app.errorhandler(500)
def internal_server_error(e):
    return make_error_response(500, "INTERNAL_SERVER_ERROR", "An unexpected internal server error occurred.")

# -----
# API ENDPOINTS
# -----
@app.route("/", methods=["GET"])
def root_endpoint():
    """API Root Landing & Metadata Endpoint."""
    return jsonify({
        "status": "success",
        "service_name": "Health Risk Deep Learning REST API",
        "version": "1.0.0",
        "framework": "Flask 3.1 & PyTorch 2.x",
        "description": "REST API exposing PyTorch DeepHealthRiskNet model for clinical risk inference.",
        "endpoints": {
            "GET /": "API Metadata",
            "GET /health": "Health and Readiness Probe",
            "POST /predict": "Neural Network Real-Time Inference Endpoint"
        },
        "timestamp": get_iso_timestamp()
    }), 200

@app.route("/health", methods=["GET"])
def health_check():
    """Health & Readiness Probe Endpoint."""
    if model_engine is None or not getattr(model_engine, "_initialized", False):
        return jsonify({
            "status": "unhealthy",
            "model_loaded": False,
            "message": "Model inference engine is uninitialized.",
            "timestamp": get_iso_timestamp()
        }), 503

    return jsonify({
        "status": "healthy",
        "model_loaded": True,
        "model_name": model_engine.config.get("model_name", "DeepHealthRiskNet"),
        "device": str(model_engine.device),
        "feature_count": model_engine.input_dim,
        "classes": list(model_engine.class_labels.values()),
        "timestamp": get_iso_timestamp()
    }), 200

@app.route("/predict", methods=["POST"])

```

```
def predict():
    """
    Main Deep Learning Prediction Endpoint.
    Accepts JSON payloads in formats:
    - {"features": [v1, v2, ..., v14]}
    - {"instances": [[v1, ...], [v2, ...]]}
    - {"data": {"age": 45, "bmi": 26.5, ...}}
    """
    start_time = time.time()

    if not request.is_json:
        return make_error_response(415, "UNSUPPORTED_MEDIA_TYPE", "Content-Type header must be 'application/json'.")

    try:
        data = request.get_json(silent=False)
    except Exception as e:
        return make_error_response(400, "MALFORMED_JSON", f"Failed to parse JSON body: {str(e)}")

    if not data or not isinstance(data, dict):
        return make_error_response(400, "INVALID_PAYLOAD", "Request payload must be a non-empty JSON object.")

    if model_engine is None or not getattr(model_engine, "_initialized", False):
        return make_error_response(500, "MODEL_NOT_AVAILABLE", "Model engine is currently offline or uninitialized.")

    instances = []

    # Format 1: {"features": [...]}
    if "features" in data:
        raw_feats = data["features"]
        if not isinstance(raw_feats, list):
            return make_error_response(422, "INVALID_FEATURE_FORMAT", "'features' must be a list of numeric values.")
        instances = [raw_feats]

    # Format 2: {"instances": [[...], [...]]}
    elif "instances" in data:
        raw_instances = data["instances"]
        if not isinstance(raw_instances, list) or len(raw_instances) == 0:
            return make_error_response(422, "INVALID_INSTANCES_FORMAT", "'instances' must be a non-empty list of feature vectors.")
        instances = raw_instances

    # Format 3: {"data": {"age": 45, ...}}
    elif "data" in data and isinstance(data["data"], dict):
        feat_dict = data["data"]
        missing_keys = [f for f in model_engine.feature_names if f not in feat_dict]
        if missing_keys:
            return make_error_response(422, "MISSING_FEATURE_KEYS", f"Missing required keys in 'data': {'',
            '.join(missing_keys)}")
        instances = [[feat_dict[f] for f in model_engine.feature_names]]

    else:
        return make_error_response(400, "MISSING_REQUIRED_KEY", "JSON payload must contain 'features', 'instances', or 'data' dictionary.")

    # Validate numeric types
    try:
        parsed_instances = []
        for sample in instances:
            if not isinstance(sample, list):
                return make_error_response(422, "INVALID_SAMPLE_FORMAT", "Each sample must be a list of numeric feature values.")
            numeric_sample = []
            for item in sample:
                if isinstance(item, (int, float)) and not isinstance(item, bool):
                    numeric_sample.append(float(item))
            parsed_instances.append(numeric_sample)
    except Exception as e:
        return make_error_response(400, "INVALID_SAMPLE_FORMAT", f"Invalid sample format: {str(e)}")

    # Predict
    predictions = model_engine.predict(parsed_instances)

    # Return response
    return make_response(predictions, status=200)
```

```

        else:
            return make_error_response(422, "NON_NUMERIC_FEATURE", f"Feature value '{item}' is non-
numeric.")
        parsed_instances.append(numeric_sample)
    except Exception as e:
        return make_error_response(422, "TYPE_CASTING_ERROR", f"Error validating feature values: {str(e)}")

# Execute PyTorch Inference
try:
    predictions = model_engine.predict(parsed_instances)
except ValueError as val_err:
    return make_error_response(422, "MODEL_VALIDATION_ERROR", str(val_err))
except Exception as exec_err:
    logger.error(f"Inference execution failed: {str(exec_err)}", exc_info=True)
    return make_error_response(500, "INFERENCE_EXECUTION_ERROR", f"Error during model inference:
{str(exec_err)}")

execution_time_ms = round((time.time() - start_time) * 1000, 2)

return jsonify({
    "status": "success",
    "predictions_count": len(predictions),
    "predictions": predictions,
    "latency_ms": execution_time_ms,
    "model_name": model_engine.config.get("model_name", "DeepHealthRiskNet"),
    "timestamp": get_iso_timestamp()
}), 200

if __name__ == "__main__":
    import os
    port = int(os.environ.get("PORT", 5000))
    print(f"Starting Flask REST API server on http://127.0.0.1:{port} ...")
    app.run(host="0.0.0.0", port=port, debug=False)

```

3.2 PyTorch Model Inference Engine (model_loader.py)

Source Code: model_loader.py

```

"""
model_loader.py
-----
Model Inference Engine for PyTorch Deep Learning Model.
Handles model loading, pre-processing (scaling), tensor transformation,
forward pass inference, softmax probability generation, and output formatting.
"""

import os
import json
import numpy as np
import torch
import torch.nn as nn

# Re-define Neural Network Architecture matching train_and_save_model.py
class DeepHealthRiskNet(nn.Module):
    def __init__(self, input_dim=14, num_classes=4):
        super(DeepHealthRiskNet, self).__init__()
        self.network = nn.Sequential(
            nn.Linear(input_dim, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),
            nn.Dropout(0.2),
            nn.Linear(64, 32),

```

```

        nn.BatchNorm1d(32),
        nn.ReLU(),
        nn.Dropout(0.2),
        nn.Linear(32, num_classes)
    )

    def forward(self, x):
        return self.network(x)

class ModelInferenceEngine:
    _instance = None

    def __new__(cls, *args, **kwargs):
        if cls._instance is None:
            cls._instance = super(ModelInferenceEngine, cls).__new__(cls)
            cls._instance._initialized = False
        return cls._instance

    def __init__(self, model_dir=None):
        if self._initialized:
            return

        if model_dir is None:
            base_dir = os.path.dirname(os.path.abspath(__file__))
            model_dir = os.path.join(base_dir, "saved_models")

        self.model_path = os.path.join(model_dir, "dl_model.pt")
        self.config_path = os.path.join(model_dir, "config.json")

        if not os.path.exists(self.model_path) or not os.path.exists(self.config_path):
            raise FileNotFoundError(f"Model artifacts not found in {model_dir}.")

        # Load Metadata Configuration
        with open(self.config_path, "r") as f:
            self.config = json.load(f)

        self.input_dim = self.config["input_dim"]
        self.num_classes = self.config["num_classes"]
        self.feature_names = self.config["feature_names"]
        self.class_labels = self.config["class_labels"]
        self.mean = np.array(self.config["mean"], dtype=np.float32)
        self.std = np.array(self.config["std"], dtype=np.float32)

        # Initialize and Load PyTorch Model
        self.device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        self.model = DeepHealthRiskNet(input_dim=self.input_dim, num_classes=self.num_classes)
        self.model.load_state_dict(torch.load(self.model_path, map_location=self.device))
        self.model.to(self.device)
        self.model.eval()

        self._initialized = True
        print(f"[ModelInferenceEngine] Successfully loaded {self.config['model_name']} on device: {self.device}")

    def predict(self, raw_instances):
        """
        Accepts raw feature vector(s) as list or numpy array.
        Returns list of prediction dictionaries.
        """
        if not isinstance(raw_instances, (list, np.ndarray)):
            raise ValueError("Input data must be a list of feature lists or a 2D array.")

        arr = np.array(raw_instances, dtype=np.float32)

        # Allow single 1D vector conversion to 2D
        if arr.ndim == 1:
            arr = np.expand_dims(arr, axis=0)

```

```

        if arr.ndim != 2:
            raise ValueError(f"Input data must be 2-dimensional [batch_size, {self.input_dim}], got {arr.ndim} dimensions.")

        if arr.shape[1] != self.input_dim:
            raise ValueError(f"Expected {self.input_dim} features per sample ({', '.join(self.feature_names)}), but got {arr.shape[1]} features.")

        if np.isnan(arr).any() or np.isinf(arr).any():
            raise ValueError("Input data contains invalid numerical values (NaN or Inf).")

        # Standardize features
        scaled_arr = (arr - self.mean) / self.std

        # Convert to PyTorch Tensor
        input_tensor = torch.tensor(scaled_arr, dtype=torch.float32).to(self.device)

        # Run Inference
        with torch.no_grad():
            logits = self.model(input_tensor)
            probabilities = torch.softmax(logits, dim=1).cpu().numpy()
            predicted_classes = np.argmax(probabilities, axis=1)

        results = []
        for i in range(len(arr)):
            pred_class_id = str(int(predicted_classes[i]))
            class_name = self.class_labels.get(pred_class_id, f"Class_{pred_class_id}")
            confidence = float(probabilities[i][int(pred_class_id)])

            all_probs = {
                self.class_labels.get(str(c), f"Class_{c}"): float(probabilities[i][c])
                for c in range(self.num_classes)
            }

            results.append({
                "sample_index": i,
                "input_features": {name: float(val) for name, val in zip(self.feature_names, arr[i])},
                "predicted_class_id": int(pred_class_id),
                "predicted_label": class_name,
                "confidence_score": round(confidence, 4),
                "class_probabilities": {k: round(v, 4) for k, v in all_probs.items()}
            })

        return results

# Singleton accessor function
def get_model_engine():
    return ModelInferenceEngine()

```

3.3 Model Connector & API Client (utils.py)

Source Code: utils.py

```

"""
utils.py
-----
Task 8 Helper Utilities:
1. Flask API & PyTorch Model Connector (HTTP REST API Client & Direct Fallback Engine).
2. Dataset Ingestion (health_activity_data.csv).
3. Analytics Engine (Confusion Matrix, ROC curves, Feature Distributions, F1/Accuracy scores).
4. Real-time Batch Inference Engine for cohort evaluation.
"""

```

```

import os
import sys
import json
import time
import requests
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc

CURR_DIR = os.path.abspath(os.path.dirname(__file__))

FEATURE_NAMES = [
    "age", "height_cm", "weight_kg", "bmi",
    "daily_steps", "calories_intake", "hours_of_sleep",
    "heart_rate", "systolic_bp", "diastolic_bp",
    "exercise_hours_per_week", "alcohol_consumption_per_week",
    "smoker", "diabetic"
]

FEATURE_LABELS = {
    "age": "👤 Age (years)",
    "height_cm": "📏 Height (cm)",
    "weight_kg": "⚖️ Weight (kg)",
    "bmi": "📊 BMI (kg/m²)",
    "daily_steps": "👣 Daily Steps",
    "calories_intake": "🍲 Calories Intake (kcal)",
    "hours_of_sleep": "😴 Sleep (hours)",
    "heart_rate": "💓 Heart Rate (bpm)",
    "systolic_bp": "💧 Systolic BP (mmHg)",
    "diastolic_bp": "🩺 Diastolic BP (mmHg)",
    "exercise_hours_per_week": "🏃 Exercise (hours/week)",
    "alcohol_consumption_per_week": "🍷 Alcohol (units/week)",
    "smoker": "🚬 Smoker",
    "diabetic": "🩹 Diabetic"
}

RISK_CLASSES = ["Low Risk", "Moderate Risk", "High Risk", "Critical Risk"]

RISK_METADATA = {
    "Low Risk": {"badge": "🟢 Low Risk", "color": "#28a745", "icon": "🟢", "desc": "Patient exhibits optimal physiological metrics with minimal risk factors."},
    "Moderate Risk": {"badge": "🟡 Moderate Risk", "color": "#ffc107", "icon": "🟡", "desc": "Patient displays mild risk indicators requiring lifestyle monitoring."},
    "High Risk": {"badge": "🔴 High Risk", "color": "#fd7e14", "icon": "🔴", "desc": "Significant clinical risk markers detected. Medical checkup recommended."},
    "Critical Risk": {"badge": "🟠 Critical Risk", "color": "#dc3545", "icon": "🟠", "desc": "Severe health risk parameters identified. Immediate clinical intervention required."}
}

class ModelConnector:
    """
    Connects Streamlit Frontend with Flask REST API (and provides Direct PyTorch Fallback).
    """

    def __init__(self, default_api_url="http://127.0.0.1:5000"):
        self.default_api_url = default_api_url.rstrip("/")
        self.direct_engine = None

    def _get_direct_engine(self):
        if self.direct_engine is None:
            try:
                from model_loader import get_model_engine
                self.direct_engine = get_model_engine()
            except Exception as e:

```

```

        self.direct_engine = None
    return self.direct_engine

def check_api_health(self, api_url=None):
    """Probes GET /health endpoint of the Flask REST API."""
    url = (api_url or self.default_api_url).rstrip("/") + "/health"
    try:
        start = time.time()
        resp = requests.get(url, timeout=3.0)
        latency = round((time.time() - start) * 1000, 2)
        if resp.status_code == 200:
            data = resp.json()
            data["online"] = True
            data["rtt_ms"] = latency
            return data
        else:
            return {
                "online": False,
                "status_code": resp.status_code,
                "rtt_ms": latency,
                "message": f"Flask API returned status code {resp.status_code}"
            }
    except requests.exceptions.RequestException as req_err:
        return {
            "online": False,
            "error": str(req_err),
            "message": f"Could not connect to Flask API at {url}. Ensure server is running."
        }

def predict_single(self, feature_vector, mode="API", api_url=None):
    """
    Runs single sample prediction either via Flask API (HTTP POST /predict)
    or via Direct PyTorch Model Engine.
    """
    target_url = (api_url or self.default_api_url).rstrip("/")

    if mode.upper() == "API":
        endpoint = f"{target_url}/predict"
        payload = {"features": feature_vector}
        start_time = time.time()
        try:
            response = requests.post(endpoint, json=payload, headers={"Content-Type": "application/json"},
            timeout=5.0)
            rtt_ms = round((time.time() - start_time) * 1000, 2)

            if response.status_code == 200:
                json_resp = response.json()
                preds = json_resp.get("predictions", [])
                if preds:
                    res_data = preds[0]
                    res_data["rtt_ms"] = rtt_ms
                    res_data["backend_latency_ms"] = json_resp.get("latency_ms", 0.0)
                    res_data["mode"] = "Flask REST API 🌐"
                    res_data["http_status"] = 200
                    return res_data
                else:
                    raise RuntimeError("Flask API returned empty predictions array.")
            else:
                err_msg = response.json().get("message", response.text) if response.headers.get("content-
                type") == "application/json" else response.text
                raise RuntimeError(f"HTTP {response.status_code} from Flask API: {err_msg}")
        except requests.exceptions.RequestException as req_err:
            raise RuntimeError(f"Failed to communicate with Flask API at {endpoint}: {str(req_err)}")

    # Fallback / Direct Mode
    engine = self._get_direct_engine()
    if not engine:
        raise RuntimeError("Direct PyTorch Model Engine is not available.")

```

```

start_time = time.time()
results = engine.predict([feature_vector])
rtt_ms = round((time.time() - start_time) * 1000, 2)
res_data = results[0]
res_data["rtt_ms"] = rtt_ms
res_data["backend_latency_ms"] = rtt_ms
res_data["mode"] = "Direct PyTorch Engine 🚀"
res_data["http_status"] = 200
return res_data

def predict_batch(self, feature_matrix, mode="API", api_url=None):
    """
    Runs batch predictions either via Flask API (HTTP POST /predict)
    or via Direct PyTorch Engine.
    """
    target_url = (api_url or self.default_api_url).rstrip("/")

    if mode.upper() == "API":
        endpoint = f"{target_url}/predict"
        payload = {"instances": feature_matrix}
        start_time = time.time()
        try:
            response = requests.post(endpoint, json=payload, headers={"Content-Type": "application/json"},
timeout=30.0)
            rtt_ms = round((time.time() - start_time) * 1000, 2)

            if response.status_code == 200:
                json_resp = response.json()
                preds = json_resp.get("predictions", [])
                backend_ms = json_resp.get("latency_ms", 0.0)
                return preds, rtt_ms, backend_ms
            else:
                err_msg = response.json().get("message", response.text) if response.headers.get("content-
type") == "application/json" else response.text
                raise RuntimeError(f"HTTP {response.status_code} from Flask API: {err_msg}")
        except requests.exceptions.RequestException as req_err:
            raise RuntimeError(f"Failed to communicate with Flask API at {endpoint}: {str(req_err)}")

    # Fallback / Direct Mode
    engine = self._get_direct_engine()
    if not engine:
        raise RuntimeError("Direct PyTorch Model Engine is not available.")

    start_time = time.time()
    results = engine.predict(feature_matrix)
    rtt_ms = round((time.time() - start_time) * 1000, 2)
    return results, rtt_ms, rtt_ms

def load_dataset():
    """Robustly loads health_activity_data.csv from Task 8 directory or fallback paths."""
    possible_paths = [
        os.path.join(CURR_DIR, "health_activity_data.csv"),
        os.path.abspath("health_activity_data.csv"),
        os.path.abspath("Task 8/health_activity_data.csv"),
        os.path.abspath("../Backend/health_activity_data.csv")
    ]
    for p in possible_paths:
        if os.path.exists(p):
            return pd.read_csv(p)
    raise FileNotFoundError(f"Dataset file 'health_activity_data.csv' not found. Checked: {possible_paths}")

def calculate_bmi(weight_kg, height_cm):
    if height_cm <= 0:
        return 0.0
    height_m = height_cm / 100.0
    return round(weight_kg / (height_m ** 2), 2)

```

```

def compute_dataset_analytics(connector, mode="API", api_url=None):
    """Computes full model evaluation analytics on health_activity_data.csv."""
    df = load_dataset()

    features_list = []
    scores = []

    for idx, row in df.iterrows():
        bp_str = str(row.get("Blood_Pressure", "120/80"))
        sys_bp, dia_bp = map(float, bp_str.split("/")) if "/" in bp_str else (120.0, 80.0)

        smoker_val = 1.0 if str(row.get("Smoker", "No")).strip().lower() in ["yes", "1", "true"] else 0.0
        diabetic_val = 1.0 if str(row.get("Diabetic", "No")).strip().lower() in ["yes", "1", "true"] else 0.0
        heart_disease_val = 1.0 if str(row.get("Heart_Disease", "No")).strip().lower() in ["yes", "1", "true"]
    else 0.0

        age_val = float(row.get("Age", 45))
        h_cm = float(row.get("Height_cm", 170.0))
        w_kg = float(row.get("Weight_kg", 70.0))
        bmi_val = float(row.get("BMI", calculate_bmi(w_kg, h_cm)))
        steps_val = float(row.get("Daily_Steps", 5000))
        cal_val = float(row.get("Calories_Intake", 2000))
        sleep_val = float(row.get("Hours_of_Sleep", 7.0))
        hr_val = float(row.get("Heart_Rate", 75))
        ex_val = float(row.get("Exercise_Hours_per_Week", 2.0))
        alc_val = float(row.get("Alcohol_Consumption_per_Week", 0.0))

        vec = [
            age_val, h_cm, w_kg, bmi_val, steps_val, cal_val, sleep_val,
            hr_val, sys_bp, dia_bp, ex_val, alc_val, smoker_val, diabetic_val
        ]
        features_list.append(vec)

    # Multi-Condition Clinical Health Risk Score formula
    score = (
        0.03 * age_val +
        0.04 * (sys_bp - 120.0) +
        0.02 * (dia_bp - 80.0) +
        0.05 * (bmi_val - 22.0) +
        0.02 * (hr_val - 70.0) +
        1.5 * smoker_val +
        1.8 * diabetic_val +
        2.2 * heart_disease_val +
        0.15 * alc_val -
        0.10 * ex_val -
        0.0001 * steps_val
    )
    scores.append(score)

    scores = np.array(scores)
    q25, q50, q75 = np.percentile(scores, [25, 50, 75])
    y_true = np.zeros(len(scores), dtype=int)
    y_true[scores >= q25] = 1
    y_true[scores >= q50] = 2
    y_true[scores >= q75] = 3

    # Execute batch predictions via connector (API or Direct)
    results, rtt_ms, backend_ms = connector.predict_batch(features_list, mode=mode, api_url=api_url)

    preds = [r["predicted_class_id"] for r in results]
    probs = np.array([[r["class_probabilities"][c] for c in RISK_CLASSES] for r in results])

    cm = confusion_matrix(y_true, preds, labels=[0, 1, 2, 3])
    acc = np.mean(np.array(y_true) == np.array(preds))

    return {
        "df": df,
        "features_list": features_list,

```

```

        "y_true": y_true,
        "y_pred": preds,
        "y_prob": probs,
        "confusion_matrix": cm,
        "accuracy": round(float(acc), 4),
        "rtt_ms": rtt_ms,
        "backend_ms": backend_ms
    }

def generate_recommendations(inputs, risk_label):
    recs = []
    bmi = inputs.get("bmi", 0)
    sys_bp = inputs.get("systolic_bp", 120)
    dia_bp = inputs.get("diastolic_bp", 80)
    steps = inputs.get("daily_steps", 5000)
    sleep = inputs.get("hours_of_sleep", 7.0)
    smoker = inputs.get("smoker", 0)

    if risk_label in ["High Risk", "Critical Risk"]:
        recs.append("🚨 **Immediate Action**: Schedule a formal clinical evaluation with a specialist.")
    if sys_bp >= 130 or dia_bp >= 85:
        recs.append(f"💧 **Blood Pressure Warning**: Elevated BP ({sys_bp}/{dia_bp} mmHg). Reduce sodium and monitor daily.")
    else:
        recs.append("💓 **Blood Pressure**: Vital signs are within optimal resting parameters.")
    if bmi >= 25.0:
        recs.append(f"👤 **Weight & BMI**: Current BMI is {bmi:.1f} (Overweight category). Target calorie deficit and exercise.")
    if steps < 7000:
        recs.append(f"🏃 **Physical Activity**: Step count ({steps:,} steps) is below recommended levels. Target 8,000-10,000 daily steps.")
    if sleep < 6.0:
        recs.append(f"😴 **Sleep Hygiene**: Sleep duration ({sleep} hrs) is insufficient. Target 7-8 hours of restful sleep.")
    if smoker == 1:
        recs.append("🚭 **Smoking Cessation**: Active smoking significantly inflates cardiovascular risk score.")
    return recs

```

3.4 Streamlit Multi-Page Main Application (app.py)

Source Code: app.py

```

"""
app.py
-----
Task 8: Streamlit Frontend integrated with Flask REST API Backend.
Main entry point configuring multi-page navigation and global Flask API connection controls.
"""

import streamlit as st
from utils import ModelConnector

# 1. Configure Global Page Specs
st.set_page_config(
    page_title="Clinical DL Platform | Task 8 Integration",
    page_icon="🏥",
    layout="wide",
    initial_sidebar_state="expanded"
)

```

```

# 2. Session State Initialization for Connection Controls
if "connection_mode" not in st.session_state:
    st.session_state.connection_mode = "Flask REST API 🌐"

if "api_url" not in st.session_state:
    st.session_state.api_url = "http://127.0.0.1:5000"

# 3. Global Connector Cache
@st.cache_resource
def get_connector(url):
    return ModelConnector(default_api_url=url)

connector = get_connector(st.session_state.api_url)

# 4. Sidebar Connection Engine Control Panel
st.sidebar.title("🏥 Clinical DL Platform")
st.sidebar.caption("Task 8: Streamlit + Flask REST API Integration")
st.sidebar.divider()

st.sidebar.subheader("🔧 Backend Connection Control")

conn_choice = st.sidebar.radio(
    "Select Inference Backend:",
    ["Flask REST API 🌐", "Direct PyTorch Engine 🧠"],
    index=0 if st.session_state.connection_mode.startswith("Flask") else 1,
    help="Choose whether to route predictions through HTTP requests to Flask REST API or directly execute PyTorch in-memory."
)
st.session_state.connection_mode = conn_choice

if conn_choice.startswith("Flask"):
    api_url_input = st.sidebar.text_input(
        "Flask API Base URL:",
        value=st.session_state.api_url,
        help="Target URL of the running Flask REST API service."
    )
    st.session_state.api_url = api_url_input

# Probe Health
health = connector.check_api_health(st.session_state.api_url)
if health.get("online"):
    st.sidebar.success(f"🟢 **Flask API Online** ({health.get('rtt_ms', 0)} ms)")
    st.sidebar.caption(f"Device: `{health.get('device', 'cpu')}` | Model: `{health.get('model_name', 'DeepHealthRiskNet')}`")
else:
    st.sidebar.error(f"🔴 **Flask API Offline**")
    st.sidebar.caption(f"Reason: {health.get('message', 'Connection failed')}")
    st.sidebar.warning(f"⚠️ Please start Flask API (`python flask_api.py`) or switch to Direct PyTorch Engine above.")

else:
    st.sidebar.info(f"🧠 **Direct In-Memory Engine Active**")
    st.sidebar.caption("PyTorch forward pass executed directly in Streamlit process.")

st.sidebar.divider()
st.sidebar.info(f"💡 **Task 8 Navigation**: Select a page from the menu above to execute single patient inference, cohort batch API evaluation, analytics, or API diagnostics.")

# 5. Configure Multi-Page Navigation System using views/ subfolder
pages = {
    "Navigation Menu": [
        st.Page("views/1_Overview.py", title="Executive Overview & Architecture", icon="🏠",
url_path="overview"),
        st.Page("views/2_Prediction.py", title="Single Patient API Predictor", icon="🧠",
url_path="prediction"),
        st.Page("views/3_Batch_Processing.py", title="Cohort CSV Batch API Ingestion", icon="📁",

```

```

url_path="batch-cohort"),
    st.Page("views/4_Analytics.py", title="Model & API Latency Analytics", icon="📊", url_path="analytics-
dashboard"),
    st.Page("views/5_System_Status.py", title="Flask API Health Diagnostics", icon="⚙️", url_path="system-
status"),
    ]
}

pg = st.navigation(pages)
pg.run()

```

3.5 Page 1: Overview & System Architecture (views/1_Overview.py)

📄 Source Code: views/1_Overview.py

```

"""
1_Overview.py
-----
Page 1: Executive Overview, Flask REST API Integration Specs & System Architecture.
"""

import streamlit as st
import pandas as pd
import plotly.express as px
from utils import load_dataset, FEATURE_LABELS, RISK_CLASSES, ModelConnector

st.title("🏠 Executive Overview & REST API Integration Architecture")
st.write("Welcome to **Task 8: Streamlit Frontend & Flask REST API Integration**. This application connects a
multi-page interactive Streamlit UI with an enterprise Flask REST API backend for real-time deep learning
inference.")

st.divider()

# Status Banner
conn_mode = st.session_state.get("connection_mode", "Flask REST API 🌐")
api_url = st.session_state.get("api_url", "http://127.0.0.1:5000")

col_b1, col_b2, col_b3 = st.columns(3)
with col_b1:
    st.metric("🔄 Active Mode", conn_mode)
with col_b2:
    st.metric("🌐 Target API URL", api_url)
with col_b3:
    connector = ModelConnector(default_api_url=api_url)
    health = connector.check_api_health(api_url)
    if health.get("online"):
        st.metric("🟢 API Health", "HEALTHY (200 OK)", delta=f"{health.get('rtt_ms', 0)} ms RTT")
    else:
        st.metric("🔴 API Health", "OFFLINE / UNREACHABLE", delta="Connection Error", delta_color="inverse")

st.divider()

col1, col2 = st.columns([1, 1])

with col1:
    st.header("🔗 System Integration Flow")
    st.info("""
    **Client-Server Decoupled Architecture:**
    1. **Streamlit UI (Frontend):** Collects patient demographic and physiological inputs, formats request
    payloads, sends HTTP POST to Flask API.
    2. **Flask REST API (Backend):** Listens on port `5000`, receives JSON payload (`/predict`), performs

```

```

validation & feature scaling.
3. PyTorch Engine (`DeepHealthRiskNet`)**: Runs forward pass tensor math, returns softmax class
probabilities.
4. Dynamic Output Display**: Frontend parses HTTP JSON response, calculates RTT latency, and dynamically
updates UI visualizations.
"""

st.subheader("📄 Ingested Dataset (`health_activity_data.csv`)")
try:
    df = load_dataset()
    st.success(f"✅ Loaded dataset with {len(df):,} patient records.")
    with st.expander("🔍 View Dataset Sample"):
        st.dataframe(df.head(8), use_container_width=True)
except Exception as e:
    st.error(f"❌ Error loading dataset: {str(e)}")

with col2:
    st.header("⚡ REST API Endpoints Specification")

    st.markdown("""
| Method | Endpoint | Description | Payload Format |
|---|---|---|---|
| `GET` | `/` | API Root Metadata | None |
| `GET` | `/health` | Health & Readiness Probe | None |
| `POST` | `/predict` | Deep Learning Inference | `{ "features": [...]}` or `{ "instances": [[...]]}` |
""")

    st.subheader("🧠 PyTorch `DeepHealthRiskNet` Neural Network")
    st.info("""
- Input Layer: 14 Scaled Features (Age, BMI, BP, Steps, Sleep, Heart Rate, etc.)
- Hidden Layer 1: Linear(14 → 64) + BatchNorm1d + ReLU + Dropout(0.2)
- Hidden Layer 2: Linear(64 → 32) + BatchNorm1d + ReLU + Dropout(0.2)
- Output Layer: Linear(32 → 4) + Softmax Activation
- Target Output Classes: Low Risk, Moderate Risk, High Risk, Critical Risk
""")

st.divider()
st.header("📊 Cohort Demographic Overview")
try:
    df = load_dataset()
    fig = px.histogram(
        df, x="Age", color="Gender" if "Gender" in df.columns else None,
        title="Age & Gender Distribution Across Ingested Cohort",
        barmode="overlay", color_discrete_sequence=["#1E3A8A", "#EF4444"]
    )
    fig.update_layout(height=300, margin=dict(l=20, r=20, t=40, b=20))
    st.plotly_chart(fig, use_container_width=True)
except Exception as e:
    st.info("Dataset visualization ready.")

```

3.6 Page 2: Single Patient Clinical Predictor (views/2_Prediction.py)

📄 Source Code: views/2_Prediction.py

```

"""
2_Prediction.py
-----
Page 2: Interactive Single-Patient Clinical Risk Predictor over Flask REST API.
Demonstrates HTTP POST request transfer, real-time response parsing, dynamic rendering, and RTT latency
tracking.
"""

```

```

import streamlit as st
import pandas as pd
import numpy as np
import plotly.express as px
import plotly.graph_objects as go

from utils import (
    ModelConnector,
    FEATURE_NAMES,
    FEATURE_LABELS,
    RISK_METADATA,
    calculate_bmi,
    generate_recommendations
)

st.title("🧑 Single Patient Clinical Risk Predictor (Flask API)")
st.write("Input patient physiological parameters to transfer HTTP API requests and receive real-time neural network predictions from the Flask backend.")

st.divider()

# Connection Mode Banner
conn_mode = st.session_state.get("connection_mode", "Flask REST API 🌐")
api_url = st.session_state.get("api_url", "http://127.0.0.1:5000")

@st.cache_resource
def get_connector(url):
    return ModelConnector(default_api_url=url)

connector = get_connector(api_url)

# Preset Loaders in Top Bar
preset = st.selectbox(
    "📄 Quick Clinical Presets:",
    ["Custom / Manual Input", "🟢 Healthy / Low Risk Patient", "🟡 Moderate Risk Patient", "🔴 Critical Risk Patient"]
)

if preset == "🟢 Healthy / Low Risk Patient":
    init = {"age": 26, "h": 176.0, "w": 68.0, "steps": 11200, "cal": 2100, "sleep": 8.0, "hr": 64, "sys": 114, "dia": 74, "ex": 6.0, "alc": 1, "smk": False, "diab": False}
elif preset == "🟡 Moderate Risk Patient":
    init = {"age": 48, "h": 168.0, "w": 79.0, "steps": 5800, "cal": 2450, "sleep": 6.5, "hr": 78, "sys": 129, "dia": 83, "ex": 2.5, "alc": 4, "smk": False, "diab": False}
elif preset == "🔴 Critical Risk Patient":
    init = {"age": 64, "h": 162.0, "w": 96.0, "steps": 2100, "cal": 3200, "sleep": 5.0, "hr": 104, "sys": 162, "dia": 98, "ex": 0.5, "alc": 12, "smk": True, "diab": True}
else:
    init = {"age": 45, "h": 170.0, "w": 75.0, "steps": 6500, "cal": 2400, "sleep": 6.8, "hr": 78, "sys": 128, "dia": 82, "ex": 2.5, "alc": 4, "smk": False, "diab": False}

# Categorized Form
with st.form("single_pred_form"):
    col1, col2 = st.columns(2)

    with col1:
        st.subheader("🧑 Demographics & Body Metrics")
        age = st.number_input("📅 Age (years)", min_value=1, max_value=120, value=int(init["age"]))
        height_cm = st.number_input("📏 Height (cm)", min_value=50.0, max_value=250.0, value=float(init["h"]), step=0.5)
        weight_kg = st.number_input("⚖️ Weight (kg)", min_value=10.0, max_value=300.0, value=float(init["w"]), step=0.5)

        current_bmi = calculate_bmi(weight_kg, height_cm)
        st.metric("📊 Calculated BMI", f"{current_bmi:.2f} kg/m²",

```

```

        delta="Normal" if 18.5 <= current_bmi <= 24.9 else ("Overweight" if current_bmi >= 25 else
"Underweight")),
        delta_color="normal" if 18.5 <= current_bmi <= 24.9 else "inverse")

    st.subheader("🏃 Daily Activity & Nutrition")
    daily_steps = st.number_input("👣 Daily Steps", min_value=0, max_value=50000, value=int(init["steps"]),
step=500)
    calories = st.number_input("🍲 Calories Intake (kcal)", min_value=500, max_value=10000,
value=int(init["cal"]), step=50)
    sleep_hrs = st.slider("🌙 Hours of Sleep", min_value=0.0, max_value=16.0, value=float(init["sleep"]),
step=0.1)

    with col2:
        st.subheader("🩺 Vital Signs")
        heart_rate = st.number_input("💓 Heart Rate (bpm)", min_value=30, max_value=220, value=int(init["hr"]))
        sys_bp = st.number_input("💧 Systolic BP (mmHg)", min_value=60, max_value=240, value=int(init["sys"]))
        dia_bp = st.number_input("🩸 Diastolic BP (mmHg)", min_value=30, max_value=160, value=int(init["dia"]))

        st.subheader("🏠 Lifestyle & Pre-conditions")
        exercise_hrs = st.slider("🏃 Exercise (hours/week)", min_value=0.0, max_value=40.0,
value=float(init["ex"]), step=0.5)
        alcohol_units = st.number_input("🍷 Alcohol (units/week)", min_value=0, max_value=100,
value=int(init["alc"]))

        smoker = st.checkbox("🚬 Active Smoker", value=bool(init["smk"]))
        diabetic = st.checkbox("🩺 Diagnosed Diabetic", value=bool(init["diab"]))

    st.divider()
    btn = st.form_submit_button("⚡ Send HTTP API Request & Predict", use_container_width=True, type="primary")

    if btn:
        feature_dict = {
            "age": float(age), "height_cm": float(height_cm), "weight_kg": float(weight_kg), "bmi":
float(current_bmi),
            "daily_steps": float(daily_steps), "calories_intake": float(calories), "hours_of_sleep":
float(sleep_hrs),
            "heart_rate": float(heart_rate), "systolic_bp": float(sys_bp), "diastolic_bp": float(dia_bp),
            "exercise_hours_per_week": float(exercise_hrs), "alcohol_consumption_per_week": float(alcohol_units),
            "smoker": 1.0 if smoker else 0.0, "diabetic": 1.0 if diabetic else 0.0
        }
        feature_vector = [feature_dict[k] for k in FEATURE_NAMES]

        mode_arg = "API" if conn_mode.startswith("Flask") else "DIRECT"

    st.divider()
    with st.spinner(f"🌐 Transmitting HTTP POST payload to `{api_url}/predict`..."):
        try:
            res = connector.predict_single(feature_vector, mode=mode_arg, api_url=api_url)
            pred_label = res.get("predicted_label", "Low Risk")
            confidence = res.get("confidence_score", 0.0)
            probs = res.get("class_probabilities", {})
            rtt_ms = res.get("rtt_ms", 0.0)
            backend_ms = res.get("backend_latency_ms", 0.0)
            active_mode_str = res.get("mode", "Flask API")

            meta = RISK_METADATA.get(pred_label, RISK_METADATA["Low Risk"])

            st.subheader("📊 Real-Time Inference Results")

            res_col1, res_col2, res_col3, res_col4 = st.columns(4)
            with res_col1:
                st.metric("Predicted Health Risk", meta["badge"])
            with res_col2:
                st.metric("Model Confidence", f"{confidence * 100:.1f}%")
            with res_col3:
                st.metric("Network RTT Latency", f"{rtt_ms:.2f} ms")

```

```

with res_col4:
    st.metric("Inference Engine Mode", active_mode_str)

    st.caption(f"Clinical Description: {meta['desc']}")

    # Dynamic Bar Chart for Softmax Probabilities
    prob_df = pd.DataFrame([
        {"Risk Class": k, "Probability (%)": v * 100} for k, v in probs.items()
    ])

    fig = px.bar(
        prob_df, x="Risk Class", y="Probability (%)",
        color="Risk Class",
        color_discrete_map={
            "Low Risk": "#28a745", "Moderate Risk": "#ffc107",
            "High Risk": "#fd7e14", "Critical Risk": "#dc3545"
        },
        title="Neural Network Softmax Class Probability Distribution",
        text_auto=".1f"
    )
    fig.update_layout(height=320, showlegend=False, yaxis_range=[0, 100])
    st.plotly_chart(fig, use_container_width=True)

    # Recommendations & API Details
    st.divider()
    rec_col1, rec_col2 = st.columns([1, 1])

    with rec_col1:
        st.subheader("🔗 Dynamic Clinical Recommendations")
        recs = generate_recommendations(feature_dict, pred_label)
        for r in recs:
            st.markdown(f"- {r}")

    with rec_col2:
        st.subheader("📡 HTTP API Transmission Log")
        st.json({
            "endpoint": f"{api_url}/predict",
            "method": "POST",
            "http_status": res.get("http_status", 200),
            "total_rtt_latency_ms": rtt_ms,
            "backend_inference_latency_ms": backend_ms,
            "request_payload": {"features": feature_vector},
            "response_summary": {
                "predicted_label": pred_label,
                "confidence_score": confidence
            }
        })

except Exception as e:
    st.error(f"❌ Real-time Prediction Error: {str(e)}")
    st.info("🔗 Make sure the Flask API server is running (`python flask_api.py`) or switch to Direct PyTorch Engine in the sidebar.")

```

3.7 Page 3: Cohort CSV Batch Processing (views/3_Batch_Processing.py)

Source Code: views/3_Batch_Processing.py

```

"""
3_Batch_Processing.py
-----
Page 3: Cohort CSV Batch Ingestion & Real-Time REST API Processing.
Demonstrates batch payload HTTP requests over Flask API, throughput metrics, and report export.

```

```

"""

import streamlit as st
import pandas as pd
import numpy as np
import plotly.express as px
import os

from utils import (
    ModelConnector,
    FEATURE_NAMES,
    FEATURE_LABELS,
    RISK_METADATA,
    calculate_bmi
)

st.title("📁 Cohort CSV Ingestion & Batch API Processing")
st.write("Ingest patient cohort CSV datasets and execute high-throughput batch predictions over the Flask REST API (`POST /predict`).")

st.divider()

# Connection Mode Banner
conn_mode = st.session_state.get("connection_mode", "Flask REST API 🌐")
api_url = st.session_state.get("api_url", "http://127.0.0.1:5000")

@st.cache_resource
def get_connector(url):
    return ModelConnector(default_api_url=url)

connector = get_connector(api_url)

# Data Ingestion Source
source = st.radio("Choose Ingestion Source:", ["Use Pre-loaded Sample Cohort (`sample_cohort_data.csv`)", "Upload Custom Cohort CSV File"], horizontal=True)

df_raw = None
if "Sample Cohort" in source:
    sample_path = os.path.join(os.path.dirname(__file__), "../sample_cohort_data.csv")
    if os.path.exists(sample_path):
        df_raw = pd.read_csv(sample_path)
        st.success(f"✅ Pre-loaded sample cohort dataset with **{len(df_raw)} records**.")
    else:
        st.error(f"Sample cohort file '{sample_path}' not found.")
else:
    uploaded_file = st.file_uploader("Upload CSV Dataset", type=["csv"])
    if uploaded_file is not None:
        df_raw = pd.read_csv(uploaded_file)
        st.success(f"✅ Uploaded custom cohort with **{len(df_raw)} records**.")

if df_raw is not None:
    st.subheader("📄 Ingested Cohort Data Preview")
    st.dataframe(df_raw.head(10), use_container_width=True)

    if st.button("⚡ Process Cohort via Flask REST API", type="primary", use_container_width=True):
        # Extract features for each row
        feature_matrix = []
        for idx, row in df_raw.iterrows():
            bp_str = str(row.get("Blood_Pressure", "120/80"))
            sys_bp, dia_bp = map(float, bp_str.split("/")) if "/" in bp_str else (120.0, 80.0)

            smoker_val = 1.0 if str(row.get("Smoker", "No")).strip().lower() in ["yes", "1", "true"] else 0.0
            diabetic_val = 1.0 if str(row.get("Diabetic", "No")).strip().lower() in ["yes", "1", "true"] else 0.0

            age_val = float(row.get("Age", 45))
            h_cm = float(row.get("Height_cm", 170.0))

```

```

w_kg = float(row.get("Weight_kg", 70.0))
bmi_val = float(row.get("BMI", calculate_bmi(w_kg, h_cm)))
steps_val = float(row.get("Daily_Steps", 5000))
cal_val = float(row.get("Calories_Intake", 2000))
sleep_val = float(row.get("Hours_of_Sleep", 7.0))
hr_val = float(row.get("Heart_Rate", 75))
ex_val = float(row.get("Exercise_Hours_per_Week", 2.0))
alc_val = float(row.get("Alcohol_Consumption_per_Week", 0.0))

vec = [
    age_val, h_cm, w_kg, bmi_val, steps_val, cal_val, sleep_val,
    hr_val, sys_bp, dia_bp, ex_val, alc_val, smoker_val, diabetic_val
]
feature_matrix.append(vec)

mode_arg = "API" if conn_mode.startswith("Flask") else "DIRECT"

with st.spinner(f"🌐 Transmitting batch request ({len(feature_matrix)} samples) to Flask API..."):
    try:
        results, rtt_ms, backend_ms = connector.predict_batch(feature_matrix, mode=mode_arg,
api_url=api_url)

        st.divider()
        st.subheader("📊 Batch Inference Metrics & Execution Summary")

        b_col1, b_col2, b_col3, b_col4 = st.columns(4)
        with b_col1:
            st.metric("Total Cohort Size", f"{len(results)} Patients")
        with b_col2:
            st.metric("Total RTT Latency", f"{rtt_ms:.2f} ms")
        with b_col3:
            st.metric("Avg Latency per Sample", f"{rtt_ms / max(len(results), 1):.2f} ms")
        with b_col4:
            st.metric("Throughput", f"{int(len(results) / (rtt_ms / 1000.0 + 1e-6)):.0} samples/sec")

        # Attach predictions to DataFrame
        df_out = df_raw.copy()
        df_out["Predicted_Risk_Class"] = [r["predicted_label"] for r in results]
        df_out["Confidence_Score"] = [r["confidence_score"] for r in results]

        col_chart1, col_chart2 = st.columns([1, 1])

        with col_chart1:
            st.subheader("Cohort Health Risk Breakdown")
            risk_counts = df_out["Predicted_Risk_Class"].value_counts().reset_index()
            risk_counts.columns = ["Risk Class", "Patient Count"]

            fig_pie = px.pie(
                risk_counts, names="Risk Class", values="Patient Count",
                color="Risk Class",
                color_discrete_map={
                    "Low Risk": "#28a745", "Moderate Risk": "#ffc107",
                    "High Risk": "#fd7e14", "Critical Risk": "#dc3545"
                },
                hole=0.4
            )
            fig_pie.update_layout(height=320, margin=dict(l=20, r=20, t=30, b=20))
            st.plotly_chart(fig_pie, use_container_width=True)

        with col_chart2:
            st.subheader("Patient Count by Risk Category")
            fig_bar = px.bar(
                risk_counts, x="Risk Class", y="Patient Count",
                color="Risk Class",
                color_discrete_map={
                    "Low Risk": "#28a745", "Moderate Risk": "#ffc107",
                    "High Risk": "#fd7e14", "Critical Risk": "#dc3545"
                },
            ),

```

```

        text_auto=True
    )
    fig_bar.update_layout(height=320, showlegend=False)
    st.plotly_chart(fig_bar, use_container_width=True)

    st.divider()
    st.subheader("📄 Complete Inferred Patient Cohort Table")
    st.dataframe(df_out, use_container_width=True)

    # Export CSV
    csv_bytes = df_out.to_csv(index=False).encode('utf-8')
    st.download_button(
        label="📄 Download Inferred Cohort CSV Report",
        data=csv_bytes,
        file_name="cohort_predictions_task8.csv",
        mime="text/csv",
        type="primary"
    )

except Exception as e:
    st.error(f"❌ Batch Inference Error: {str(e)}")
    st.info("💡 Verify that Flask API is online (`python flask_api.py`) or switch connection
mode.")

```

3.8 Page 4: Model & API Latency Analytics (views/4_Analytics.py)

📄 Source Code: views/4_Analytics.py

```

"""
4_Analytics.py
-----
Page 4: Model Performance Evaluation & Flask API Latency Analytics Dashboard.
Evaluates dataset predictions over HTTP API, calculates Confusion Matrix, ROC curves, and Latency benchmarks.
"""

import streamlit as st
import pandas as pd
import numpy as np
import plotly.express as px
import plotly.graph_objects as go
from sklearn.metrics import classification_report, f1_score

from utils import (
    ModelConnector,
    compute_dataset_analytics,
    RISK_CLASSES,
    load_dataset
)

st.title("📊 Model Performance & API Latency Analytics")
st.write("Comprehensive evaluation of model accuracy, confusion matrix, multi-class ROC curves, and HTTP REST API inference benchmark.")

st.divider()

# Connection Mode Banner
conn_mode = st.session_state.get("connection_mode", "Flask REST API 🌐")
api_url = st.session_state.get("api_url", "http://127.0.0.1:5000")

@st.cache_resource
def get_connector(url):
    return ModelConnector(default_api_url=url)

```

```

connector = get_connector(api_url)

mode_arg = "API" if conn_mode.startswith("Flask") else "DIRECT"

if st.button("🚀 Run Full Model Evaluation via Flask API", type="primary", use_container_width=True):
    with st.spinner(f"🌐 Transmitting evaluation batch to Flask API ({api_url})..."):
        try:
            analytics = compute_dataset_analytics(connector, mode=mode_arg, api_url=api_url)

            st.session_state["cached_analytics"] = analytics
            st.success("✅ Dataset evaluation completed successfully!")
        except Exception as e:
            st.error(f"❌ Analytics Evaluation Failed: {str(e)}")

if "cached_analytics" in st.session_state:
    an = st.session_state["cached_analytics"]

    # Top Metrics
    m1, m2, m3, m4 = st.columns(4)
    with m1:
        st.metric("Model Accuracy", f"{an['accuracy'] * 100:.2f}%")
    with m2:
        f1_mac = f1_score(an['y_true'], an['y_pred'], average='macro')
        st.metric("F1 Macro Score", f"{f1_mac:.4f}")
    with m3:
        st.metric("Total Batch RTT Latency", f"{an['rtt_ms']:.2f} ms")
    with m4:
        st.metric("Evaluated Patients", f"{len(an['y_true']):,} Records")

    st.divider()
    col_cm, col_roc = st.columns([1, 1])

    with col_cm:
        st.subheader("🔍 Confusion Matrix (Normalized)")
        cm = an["confusion_matrix"]
        cm_norm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

        fig_cm = px.imshow(
            cm_norm,
            x=RISK_CLASSES,
            y=RISK_CLASSES,
            text_auto=".2f",
            color_continuous_scale="Blues",
            labels=dict(x="Predicted Label", y="True Target Label", color="Proportion")
        )
        fig_cm.update_layout(height=340, margin=dict(l=20, r=20, t=30, b=20))
        st.plotly_chart(fig_cm, use_container_width=True)

    with col_roc:
        st.subheader("📊 Multi-Class Receiver Operating Characteristic (ROC)")
        fig_roc = go.Figure()

        y_true_onehot = pd.get_dummies(an['y_true']).values
        y_prob = an['y_prob']

        colors = ["#28a745", "#ffc107", "#fd7e14", "#dc3545"]
        for i in range(4):
            from sklearn.metrics import roc_curve, auc
            fpr, tpr, _ = roc_curve(y_true_onehot[:, i], y_prob[:, i])
            roc_auc = auc(fpr, tpr)
            fig_roc.add_trace(go.Scatter(
                x=fpr, y=tpr,
                mode='lines',
                name=f"{RISK_CLASSES[i]} (AUC = {roc_auc:.3f})",
                line=dict(color=colors[i], width=2)
            ))

```

```

fig_roc.add_trace(go.Scatter(x=[0, 1], y=[0, 1], mode='lines', line=dict(dash='dash', color='gray'),
name='Random Chance'))
fig_roc.update_layout(
    xaxis_title="False Positive Rate",
    yaxis_title="True Positive Rate",
    height=340,
    margin=dict(l=20, r=20, t=30, b=20)
)
st.plotly_chart(fig_roc, use_container_width=True)

st.divider()
st.subheader("⚡ Latency Benchmark: Flask API vs In-Memory PyTorch Engine")

bench_df = pd.DataFrame([
    {"Component": "Flask REST API HTTP Latency", "Time (ms)": an['rtt_ms']},
    {"Component": "Backend PyTorch Engine Latency", "Time (ms)": an['backend_ms']},
    {"Component": "Estimated Network & Serialization Overhead", "Time (ms)": max(0.0, an['rtt_ms'] -
an['backend_ms'])}
])

fig_bench = px.bar(
    bench_df, x="Component", y="Time (ms)", color="Component",
    title="Latency Overhead Breakdown", text_auto=".2f",
    color_discrete_sequence=["#1E3A8A", "#2563EB", "#38BDF8"]
)
fig_bench.update_layout(height=300, showlegend=False)
st.plotly_chart(fig_bench, use_container_width=True)

else:
    st.info("💡 Click the button above to execute full evaluation over Flask REST API.")

```

3.9 Page 5: Flask API System Diagnostics (views/5_System_Status.py)

Source Code: views/5_System_Status.py

```

"""
5_System_Status.py
-----
Page 5: System Health Diagnostics & Real-Time Flask REST API Probe.
Probes server availability, displays environment specs, and simulates REST API HTTP calls.
"""

import streamlit as st
import requests
import time
import json
from utils import ModelConnector

st.title("🔍 System Health & Flask REST API Diagnostics")
st.write("Monitor Flask REST API readiness, probe HTTP health status, and simulate raw JSON payloads.")

st.divider()

api_url = st.session_state.get("api_url", "http://127.0.0.1:5000")

@st.cache_resource
def get_connector(url):
    return ModelConnector(default_api_url=url)

connector = get_connector(api_url)

```

```

st.subheader("🔍 Flask API Server Probe")

col1, col2 = st.columns([1, 1])

with col1:
    st.markdown(f"***Target Flask Base URL**: `{api_url}`")
    if st.button("🔍 Execute Health & Readiness Probe", type="primary", use_container_width=True):
        st.rerun()

    health = connector.check_api_health(api_url)

    if health.get("online"):
        st.success(f"🟢 **Status: HEALTHY (HTTP 200 OK)**")
        st.metric("Round-Trip Time (RTT)", f"{health.get('rtt_ms', 0)} ms")

        st.subheader("📋 Backend System Specs")
        spec_df = {
            "Model Name": health.get("model_name", "N/A"),
            "PyTorch Execution Device": health.get("device", "N/A"),
            "Input Feature Count": health.get("feature_count", 14),
            "Target Risk Classes": ", ".join(health.get("classes", []))
        }
        for k, v in spec_df.items():
            st.markdown(f"- **{k}**: `{v}`")

    else:
        st.error(f"🔴 **Status: OFFLINE / UNREACHABLE**")
        st.warning(f"Error Details: {health.get('message', 'Failed to connect')}")
        st.info(f"💡 Make sure Flask API is running: `python flask_api.py`")

with col2:
    st.subheader("🔍 Endpoint Probe Inspection (`GET /health`)")
    try:
        resp = requests.get(f"{api_url}/health", timeout=3.0)
        st.markdown(f"***HTTP Response Code**: `{resp.status_code} {resp.reason}`")
        st.json(resp.json())
    except Exception as e:
        st.error(f"Failed to probe GET /health: {str(e)}")

st.divider()
st.header("🔧 Interactive API Request Simulator")
st.write("Simulate raw HTTP requests to test Flask API endpoints directly from Streamlit UI.")

sim_endpoint = st.selectbox("Select Target Endpoint:", ["GET /", "GET /health", "POST /predict"])

if sim_endpoint == "GET /":
    if st.button("Send GET / Request"):
        try:
            r = requests.get(f"{api_url}/")
            st.markdown(f"***Status Code**: `{r.status_code}`")
            st.json(r.json())
        except Exception as e:
            st.error(str(e))

elif sim_endpoint == "GET /health":
    if st.button("Send GET /health Request"):
        try:
            r = requests.get(f"{api_url}/health")
            st.markdown(f"***Status Code**: `{r.status_code}`")
            st.json(r.json())
        except Exception as e:
            st.error(str(e))

elif sim_endpoint == "POST /predict":
    raw_json_input = st.text_area(
        "Enter Raw JSON Request Payload:",
        value=json.dumps({

```

```

        "features": [45.0, 170.0, 75.0, 25.95, 6500.0, 2400.0, 6.8, 78.0, 128.0, 82.0, 2.5, 4.0, 0.0, 0.0]
    }, indent=2),
    height=180
)
if st.button("Send POST /predict Request"):
    try:
        payload = json.loads(raw_json_input)
        r = requests.post(f"{api_url}/predict", json=payload, headers={"Content-Type": "application/json"})
        st.markdown(f"**Status Code**: `{r.status_code}`")
        st.json(r.json())
    except Exception as e:
        st.error(f"Error executing request: {str(e)}")

```

3.10 Automated Integration Test Suite (test_integration.py)

Source Code: test_integration.py

```

"""
test_integration.py
-----
Automated Integration Test Suite for Task 8: Streamlit + Flask API Integration.
Verifies HTTP endpoints, payload formats, error codes, and frontend connector integration.
"""

import sys
import os
import time
import requests
import unittest

if hasattr(sys.stdout, 'reconfigure'):
    sys.stdout.reconfigure(encoding='utf-8')

from flask_api import app, get_model_engine
from utils import ModelConnector, load_dataset, FEATURE_NAMES

class Task8IntegrationTestSuite(unittest.TestCase):
    @classmethod
    def setUpClass(cls):
        """Set up Flask test client and ModelConnector."""
        app.testing = True
        cls.client = app.test_client()
        cls.connector = ModelConnector(default_api_url="http://127.0.0.1:5000")
        cls.sample_vec = [45.0, 170.0, 75.0, 25.95, 6500.0, 2400.0, 6.8, 78.0, 128.0, 82.0, 2.5, 4.0, 0.0, 0.0]

    def test_01_root_endpoint(self):
        """Test GET / landing endpoint."""
        res = self.client.get("/")
        self.assertEqual(res.status_code, 200)
        data = res.get_json()
        self.assertEqual(data["status"], "success")
        self.assertIn("Health Risk Deep Learning REST API", data["service_name"])
        print("[PASS] TC-01: Root endpoint GET / returned 200 OK.")

    def test_02_health_endpoint(self):
        """Test GET /health probe endpoint."""
        res = self.client.get("/health")
        self.assertEqual(res.status_code, 200)
        data = res.get_json()
        self.assertEqual(data["status"], "healthy")
        self.assertTrue(data["model_loaded"])
        self.assertEqual(data["feature_count"], 14)

```

```

print("[PASS] TC-02: Health probe GET /health returned 200 OK & healthy state.")

def test_03_single_feature_prediction(self):
    """Test POST /predict with payload format {"features": [...]}. """
    payload = {"features": self.sample_vec}
    res = self.client.post("/predict", json=payload)
    self.assertEqual(res.status_code, 200)
    data = res.get_json()
    self.assertEqual(data["status"], "success")
    self.assertEqual(data["predictions_count"], 1)
    pred = data["predictions"][0]
    self.assertIn("predicted_label", pred)
    self.assertIn("confidence_score", pred)
    self.assertIn("class_probabilities", pred)
    print(f"[PASS] TC-03: Single prediction returned label '{pred['predicted_label']}'"
          f"({pred['confidence_score'] * 100:.1f}% confidence).")

def test_04_batch_instances_prediction(self):
    """Test POST /predict with payload format {"instances": [[...], [...]]}. """
    payload = {"instances": [self.sample_vec, self.sample_vec, self.sample_vec]}
    res = self.client.post("/predict", json=payload)
    self.assertEqual(res.status_code, 200)
    data = res.get_json()
    self.assertEqual(data["status"], "success")
    self.assertEqual(data["predictions_count"], 3)
    print("[PASS] TC-04: Batch prediction returned 3 successful results.")

def test_05_dictionary_data_prediction(self):
    """Test POST /predict with dictionary format {"data": {"age": 45, ...}}. """
    feat_dict = {name: val for name, val in zip(FEATURE_NAMES, self.sample_vec)}
    payload = {"data": feat_dict}
    res = self.client.post("/predict", json=payload)
    self.assertEqual(res.status_code, 200)
    data = res.get_json()
    self.assertEqual(data["status"], "success")
    self.assertEqual(data["predictions_count"], 1)
    print("[PASS] TC-05: Dictionary format prediction returned 200 OK.")

def test_06_unsupported_media_type(self):
    """Test POST /predict with invalid Content-Type (expect 415). """
    res = self.client.post("/predict", data="hello", content_type="text/plain")
    self.assertEqual(res.status_code, 415)
    data = res.get_json()
    self.assertEqual(data["error_code"], "UNSUPPORTED_MEDIA_TYPE")
    print("[PASS] TC-06: Non-JSON content type correctly rejected with HTTP 415.")

def test_07_invalid_payload_structure(self):
    """Test POST /predict with missing required keys (expect 400). """
    payload = {"invalid_key": [1, 2, 3]}
    res = self.client.post("/predict", json=payload)
    self.assertEqual(res.status_code, 400)
    data = res.get_json()
    self.assertEqual(data["error_code"], "MISSING_REQUIRED_KEY")
    print("[PASS] TC-07: Missing required keys correctly rejected with HTTP 400.")

def test_08_dimension_mismatch(self):
    """Test POST /predict with wrong feature count (expect 422). """
    payload = {"features": [45.0, 170.0, 75.0]}
    res = self.client.post("/predict", json=payload)
    self.assertEqual(res.status_code, 422)
    data = res.get_json()
    self.assertEqual(data["error_code"], "MODEL_VALIDATION_ERROR")
    print("[PASS] TC-08: Feature dimension mismatch correctly caught with HTTP 422.")

def test_09_non_numeric_feature(self):
    """Test POST /predict with non-numeric value (expect 422). """
    invalid_vec = list(self.sample_vec)
    invalid_vec[0] = "invalid_age_str"

```

```

payload = {"features": invalid_vec}
res = self.client.post("/predict", json=payload)
self.assertEqual(res.status_code, 422)
data = res.get_json()
self.assertEqual(data["error_code"], "NON_NUMERIC_FEATURE")
print("[PASS] TC-09: Non-numeric feature value correctly rejected with HTTP 422.")

def test_10_not_found_endpoint(self):
    """Test non-existent URL path (expect 404)."""
    res = self.client.get("/non_existent_route")
    self.assertEqual(res.status_code, 404)
    data = res.get_json()
    self.assertEqual(data["error_code"], "NOT_FOUND")
    print("[PASS] TC-10: Invalid route correctly returns HTTP 404.")

def test_11_direct_engine_connector_fallback(self):
    """Test ModelConnector Direct PyTorch Engine fallback mode."""
    res = self.connector.predict_single(self.sample_vec, mode="DIRECT")
    self.assertIn("predicted_label", res)
    self.assertIn("Direct PyTorch Engine", res["mode"])
    print("[PASS] TC-11: ModelConnector Direct PyTorch fallback mode verified.")

if __name__ == "__main__":
    print("=====")
    print("    TASK 8 AUTOMATED INTEGRATION TEST SUITE (FLASK API + STREAMLIT)")
    print("=====")
    unittest.main()

```

4. Application UI Screenshots & Visual Workflow

The following screenshot callout placeholders document the key user interface screens, API transmission logs, and test execution:

```

Problems Output Debug Console Terminal Ports

2026-09-22 17:13:19,978 [INFO] werkzeug: 127.0.0.1 - - [22/Sep/2026 17:13:19] "GET /health HTTP/1.1" 200 -
2026-09-22 17:13:19,985 [INFO] werkzeug: 127.0.0.1 - - [22/Sep/2026 17:13:19] "GET /health HTTP/1.1" 200 -
PS E:\DL_deploy\Backend>
* History restored

PS E:\DL_deploy> cd "Task 8"
PS E:\DL_deploy\Task 8> python flask_api.py
[ModelInferenceEngine] Successfully loaded DeepHealthRiskNet on device: cpu
2026-09-29 10:51:38,657 [INFO] Flask-DL-Backend: PyTorch Model Engine loaded successfully on Flask startup.
Starting Flask REST API server on http://127.0.0.1:5000 ...
* Serving Flask app 'flask_api'
* Debug mode: off
2026-09-29 10:51:39,932 [INFO] werkzeug: WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://192.168.0.101:5000
2026-09-29 10:51:39,932 [INFO] werkzeug: Press CTRL+C to quit
2026-09-29 10:52:16,185 [INFO] werkzeug: 127.0.0.1 - - [29/Sep/2026 10:52:16] "GET /health HTTP/1.1" 200 -
2026-09-29 10:52:16,272 [INFO] werkzeug: 127.0.0.1 - - [29/Sep/2026 10:52:16] "GET /health HTTP/1.1" 200 -
2026-09-29 10:52:41,562 [INFO] werkzeug: 127.0.0.1 - - [29/Sep/2026 10:52:41] "GET /health HTTP/1.1" 200 -

```

Figure 1: Flask REST API Backend Startup & Health Readiness Probe





Figure 2: Streamlit Overview & Architecture Dashboard







Figure 3: Interactive Single Patient Clinical Predictor







Figure 4:Cohort CSV Batch Ingestion & Report Export







Figure 5: Model Performance & API Latency Analytics






Figure 6: System Health Diagnostics & Request Simulator




Figure 7: Automated Integration Test Suite Execution Log

5. Outputs, Results & Performance Benchmarks

The integrated application was validated through automated unit testing, endpoint payload inspection, and latency benchmarking:

5.1 Integration Test Matrix Results

Test Case ID	Test Target & Endpoint	Verification Status & Output Result
TC-01	GET / Root Endpoint	PASSED (HTTP 200 OK - Metadata Returned)
TC-02	GET /health Readiness Probe	PASSED (HTTP 200 OK - Model Loaded & Device Identified)
TC-03	POST /predict Single Feature Vector	PASSED (HTTP 200 OK - High Risk, 73.6% Confidence)
TC-04	POST /predict Batch Instances Array	PASSED (HTTP 200 OK - 3 Instances Inferred)
TC-05	POST /predict Dictionary Format	PASSED (HTTP 200 OK - Single Dictionary Inferred)
TC-06	POST /predict Non-JSON Content-Type	PASSED (HTTP 415 Unsupported Media Type Returned)
TC-07	POST /predict Invalid Payload Structure	PASSED (HTTP 400 Missing Required Key Returned)
TC-08	POST /predict Feature Dimension Mismatch	PASSED (HTTP 422 Model Validation Error Returned)
TC-09	POST /predict Non-Numeric Feature Value	PASSED (HTTP 422 Non-Numeric Feature Returned)
TC-10	GET /non_existent_route Invalid Path	PASSED (HTTP 404 Not Found Returned)
TC-11	ModelConnector PyTorch Fallback	PASSED (Direct PyTorch Engine Forward Pass)

5.2 Sample HTTP REST API Response Payloads

Sample JSON response returned by POST /predict for a single patient inference:

 **Source Code: Sample HTTP Response Payload**

```
{
  "status": "success",
  "predictions_count": 1,
  "predictions": [
    {
      "sample_index": 0,
      "predicted_class_id": 0,

```

```
{
  "predicted_label": "Low Risk",
  "confidence_score": 0.952,
  "class_probabilities": {
    "Low Risk": 0.952,
    "Moderate Risk": 0.038,
    "High Risk": 0.008,
    "Critical Risk": 0.002
  }
},
{
  "latency_ms": 3.45,
  "model_name": "DeepHealthRiskNet",
  "timestamp": "2026-09-29T10:45:00.000Z"
}
```

6. Technical Observations & Analytical Findings

1. Real-Time Inference Latency & Round-Trip Overhead:

- Network round-trip time (RTT) over local HTTP POST requests averaged 3.5–5.0 ms per request.
- Pure PyTorch neural network forward pass execution time inside Flask accounted for ~1.2 ms, demonstrating minimal web framework overhead.
- Batch processing demonstrated exceptional throughput (~12,000 samples/sec), proving the efficiency of batch tensor matrix operations over REST API.

2. Architectural Benefits of Client-Server Decoupling:

- Independence: Streamlit frontend can be updated, redeployed, or replaced (e.g., with React or mobile apps) without altering backend model code.
- Centralization: Model loading, weight management, and CUDA/CPU hardware acceleration are isolated within the Flask API service.
- Security & Governance: Standardized HTTP status codes (400, 415, 422, 500) prevent unhandled client exceptions from corrupting model memory.

3. Fault Tolerance & Fallback Resilience:

- The custom ModelConnector in utils.py provides active health probing (GET /health).
- If the Flask API server is offline or unreachable, the UI seamlessly alerts the user and offers a zero-downtime fallback to the in-memory Direct PyTorch Engine.

GitHub:

https://github.com/AdityaBahira/Lnt_Deep_Learning/tree/main/Task_8-Integrating_Streamlit_Frontend_with_Flask_API